

# Pandas -- Series and Data Frames ¶

```
In [1]: import pandas as pd
import numpy as np
```

## Series

- a one-dimensional array-like object containing a sequence of values
- associated array of data labels, called its index

```
In [2]: np.random.seed(123)
scores = np.random.randint(60, 90, 6)

a = pd.Series(scores)
a
```

```
Out[2]: 0    73
        1    62
        2    88
        3    62
        4    66
        5    77
dtype: int64
```

```
In [3]: a.values
```

```
Out[3]: array([73, 62, 88, 62, 66, 77])
```

```
In [4]: a.index
```

```
Out[4]: RangeIndex(start=0, stop=6, step=1)
```

```
In [5]: a[1]
```

```
Out[5]: 62
```

```
In [6]: a[[1, 4]]
```

```
Out[6]: 1    62
        4    66
dtype: int64
```

```
In [7]: a[::-2]
```

```
Out[7]: 5    77
        3    62
        1    62
dtype: int64
```

```
In [8]: b = pd.Series(scores, index = ['Alice', 'Bob', 'Charlie', 'Dave', 'Ed', 'Fred'])  
b
```

```
Out[8]: Alice      73  
        Bob       62  
        Charlie   88  
        Dave      62  
        Ed       66  
        Fred     77  
        dtype: int64
```

```
In [9]: b['Bob']
```

```
Out[9]: 62
```

```
In [10]: b[['Bob', 'Ed']]
```

```
Out[10]: Bob      62  
        Ed      66  
        dtype: int64
```

```
In [11]: b[::-2]
```

```
Out[11]: Fred     77  
        Dave     62  
        Bob      62  
        dtype: int64
```

```
In [12]: b[b > 70]
```

```
Out[12]: Alice     73  
        Charlie   88  
        Fred     77  
        dtype: int64
```

```
In [13]: b + 10
```

```
Out[13]: Alice     83  
        Bob       72  
        Charlie   98  
        Dave      72  
        Ed       76  
        Fred     87  
        dtype: int64
```

```
In [14]: b
```

```
Out[14]: Alice     73  
        Bob       62  
        Charlie   88  
        Dave      62  
        Ed       66  
        Fred     77  
        dtype: int64
```

```
In [15]: np.cumsum(b)
```

```
Out[15]: Alice      73
         Bob       135
         Charlie   223
         Dave      285
         Ed        351
         Fred     428
         dtype: int64
```

```
In [16]: np.average(b)
```

```
Out[16]: 71.33333333333333
```

```
In [17]: b.describe()
```

```
Out[17]: count      6.000000
         mean      71.333333
         std       10.152175
         min       62.000000
         25%       63.000000
         50%       69.500000
         75%       76.000000
         max       88.000000
         dtype: float64
```

```
In [18]: 'Charlie' in b
```

```
Out[18]: True
```

```
In [19]: 'Robert' in b
```

```
Out[19]: False
```

```
In [20]: b.index.name = 'FirstName'
         b
```

```
Out[20]: FirstName
         Alice      73
         Bob        62
         Charlie    88
         Dave        62
         Ed         66
         Fred       77
         dtype: int64
```

### Series from dictionary data

```
In [21]: c = pd.Series({'R': 60, 'Python': 75, 'Java': 50})
         c
```

```
Out[21]: R          60
         Python    75
         Java      50
         dtype: int64
```

```
In [22]: d = pd.Series({'R': 60, 'Python': 75, 'Java': 50},  
                      index=['Java', 'Python', 'R', 'C++'])  
d
```

```
Out[22]: Java      50.0  
        Python    75.0  
        R         60.0  
        C++       NaN  
        dtype: float64
```

```
In [23]: pd.isnull(d)
```

```
Out[23]: Java      False  
        Python    False  
        R         False  
        C++       True  
        dtype: bool
```

```
In [24]: pd.notnull(d)
```

```
Out[24]: Java      True  
        Python    True  
        R         True  
        C++       False  
        dtype: bool
```

```
In [25]: c + d
```

```
Out[25]: C++       NaN  
        Java      100.0  
        Python    150.0  
        R         120.0  
        dtype: float64
```

```
In [26]: (c + d).dropna()
```

```
Out[26]: Java      100.0  
        Python    150.0  
        R         120.0  
        dtype: float64
```

## DataFrame

- represents a rectangular table of data
- contains an ordered collection of columns
- each column can be a different value type
- has both a row and column index

```
In [27]: data = {'state': ['Ohio', 'Ohio', 'Ohio', 'Nevada', 'Nevada', 'Nevada'],
                'year': [2000, 2001, 2002, 2001, 2002, 2003],
                'pop': [1.5, 1.7, 3.6, 2.4, 2.9, 3.2]}

df1 = pd.DataFrame(data)
df1
```

Out[27]:

	state	year	pop
0	Ohio	2000	1.5
1	Ohio	2001	1.7
2	Ohio	2002	3.6
3	Nevada	2001	2.4
4	Nevada	2002	2.9
5	Nevada	2003	3.2

```
In [28]: df1 = pd.DataFrame(data, columns = ['year', 'state', 'pop'])
df1
```

Out[28]:

	year	state	pop
0	2000	Ohio	1.5
1	2001	Ohio	1.7
2	2002	Ohio	3.6
3	2001	Nevada	2.4
4	2002	Nevada	2.9
5	2003	Nevada	3.2

```
In [29]: df1.head()
```

Out[29]:

	year	state	pop
0	2000	Ohio	1.5
1	2001	Ohio	1.7
2	2002	Ohio	3.6
3	2001	Nevada	2.4
4	2002	Nevada	2.9

```
In [30]: df1.tail(n=3)
```

Out[30]:

	year	state	pop
3	2001	Nevada	2.4
4	2002	Nevada	2.9
5	2003	Nevada	3.2

```
In [31]: df2 = pd.DataFrame(data, columns = ['year', 'state', 'pop', 'debt'])
df2
```

```
Out[31]:
```

	year	state	pop	debt
0	2000	Ohio	1.5	NaN
1	2001	Ohio	1.7	NaN
2	2002	Ohio	3.6	NaN
3	2001	Nevada	2.4	NaN
4	2002	Nevada	2.9	NaN
5	2003	Nevada	3.2	NaN

```
In [32]: df2.columns
```

```
Out[32]: Index(['year', 'state', 'pop', 'debt'], dtype='object')
```

```
In [33]: df2.index
```

```
Out[33]: RangeIndex(start=0, stop=6, step=1)
```

## Retrieve columns

```
In [34]: df2['year']
```

```
Out[34]: 0    2000
1    2001
2    2002
3    2001
4    2002
5    2003
Name: year, dtype: int64
```

```
In [35]: df2.year
```

```
Out[35]: 0    2000
1    2001
2    2002
3    2001
4    2002
5    2003
Name: year, dtype: int64
```

```
In [36]: df2[['year', 'state']]
```

```
Out[36]:
```

	year	state
0	2000	Ohio
1	2001	Ohio
2	2002	Ohio
3	2001	Nevada
4	2002	Nevada
5	2003	Nevada

## Retrieve rows

```
In [37]: df2 = pd.DataFrame(data, columns = ['year', 'state', 'pop', 'debt'])  
df2
```

Out[37]:

	year	state	pop	debt
0	2000	Ohio	1.5	NaN
1	2001	Ohio	1.7	NaN
2	2002	Ohio	3.6	NaN
3	2001	Nevada	2.4	NaN
4	2002	Nevada	2.9	NaN
5	2003	Nevada	3.2	NaN

```
In [38]: df2.iloc[2]
```

Out[38]: year 2002  
state Ohio  
pop 3.6  
debt NaN  
Name: 2, dtype: object

```
In [39]: type(df2.iloc[2])
```

Out[39]: pandas.core.series.Series

```
In [40]: df2.iloc[[2]]
```

Out[40]:

	year	state	pop	debt
2	2002	Ohio	3.6	NaN

```
In [41]: type(df2.iloc[[2]])
```

Out[41]: pandas.core.frame.DataFrame

```
In [42]: df2.iloc[[2,5]]
```

Out[42]:

	year	state	pop	debt
2	2002	Ohio	3.6	NaN
5	2003	Nevada	3.2	NaN

```
In [43]: df2.index = ['one', 'two', 'three', 'four', 'five', 'six']
df2
```

Out[43]:

	year	state	pop	debt
<b>one</b>	2000	Ohio	1.5	NaN
<b>two</b>	2001	Ohio	1.7	NaN
<b>three</b>	2002	Ohio	3.6	NaN
<b>four</b>	2001	Nevada	2.4	NaN
<b>five</b>	2002	Nevada	2.9	NaN
<b>six</b>	2003	Nevada	3.2	NaN

```
In [44]: df2.loc['two']
```

Out[44]:

year	2001
state	Ohio
pop	1.7
debt	NaN

Name: two, dtype: object

```
In [45]: df2.loc[['two', 'five']]
```

Out[45]:

	year	state	pop	debt
<b>two</b>	2001	Ohio	1.7	NaN
<b>five</b>	2002	Nevada	2.9	NaN

```
In [ ]:
```

```
In [46]: df2['debt'] = 20.5
df2
```

Out[46]:

	year	state	pop	debt
<b>one</b>	2000	Ohio	1.5	20.5
<b>two</b>	2001	Ohio	1.7	20.5
<b>three</b>	2002	Ohio	3.6	20.5
<b>four</b>	2001	Nevada	2.4	20.5
<b>five</b>	2002	Nevada	2.9	20.5
<b>six</b>	2003	Nevada	3.2	20.5



```
In [47]: df2['debt'] = np.arange(df2.shape[0])
df2
```

Out[47]:

	year	state	pop	debt
one	2000	Ohio	1.5	0
two	2001	Ohio	1.7	1
three	2002	Ohio	3.6	2
four	2001	Nevada	2.4	3
five	2002	Nevada	2.9	4
six	2003	Nevada	3.2	5

```
In [48]: # adding a column

df2['east'] = df2.state == 'Ohio'
df2
```

Out[48]:

	year	state	pop	debt	east
one	2000	Ohio	1.5	0	True
two	2001	Ohio	1.7	1	True
three	2002	Ohio	3.6	2	True
four	2001	Nevada	2.4	3	False
five	2002	Nevada	2.9	4	False
six	2003	Nevada	3.2	5	False

```
In [49]: # deleting a column

del df2['east']
df2
```

Out[49]:

	year	state	pop	debt
one	2000	Ohio	1.5	0
two	2001	Ohio	1.7	1
three	2002	Ohio	3.6	2
four	2001	Nevada	2.4	3
five	2002	Nevada	2.9	4
six	2003	Nevada	3.2	5

```
In [50]: df2.T
```

Out[50]:

	one	two	three	four	five	six
year	2000	2001	2002	2001	2002	2003
state	Ohio	Ohio	Ohio	Nevada	Nevada	Nevada
pop	1.5	1.7	3.6	2.4	2.9	3.2
debt	0	1	2	3	4	5

```
In [51]: # nested dictionaries

pop = {'Nevada': {2001: 2.4, 2002: 2.9},
       'Ohio': {2000: 1.5, 2001: 1.7, 2002: 3.6}}

df3 = pd.DataFrame(pop)
df3
```

Out[51]:

	Nevada	Ohio
2000	NaN	1.5
2001	2.4	1.7
2002	2.9	3.6

## Reindexing

```
In [52]: df1 = pd.Series([4.5, 7.2, -5.3, 3.6], index=['d', 'b', 'a', 'c'])
df1
```

Out[52]:

d	4.5
b	7.2
a	-5.3
c	3.6

dtype: float64

```
In [53]: df2 = df1.reindex(['a', 'b', 'c', 'd', 'e'])
df2
```

Out[53]:

a	-5.3
b	7.2
c	3.6
d	4.5
e	NaN

dtype: float64

```
In [54]: df3 = pd.Series(['blue', 'purple', 'yellow'], index=[0, 2, 4])
df3
```

Out[54]:

0	blue
2	purple
4	yellow

dtype: object

```
In [55]: df3.reindex(np.arange(6))
```

Out[55]:

0	blue
1	NaN
2	purple
3	NaN
4	yellow
5	NaN

dtype: object

```
In [56]: # forward fill missing values

df3.reindex(np.arange(6), method='ffill')
```

```
Out[56]: 0      blue
1      blue
2     purple
3     purple
4     yellow
5     yellow
dtype: object
```

```
In [57]: # backward fill missing values

df3.reindex(np.arange(6), method='bfill')
```

```
Out[57]: 0      blue
1     purple
2     purple
3     yellow
4     yellow
5         NaN
dtype: object
```

```
In [58]: df4 = pd.DataFrame(np.arange(9).reshape((3, 3)),
                             index=['a', 'c', 'd'],
                             columns=['Ohio', 'Texas', 'California'])
df4
```

```
Out[58]:
```

	Ohio	Texas	California
a	0	1	2
c	3	4	5
d	6	7	8

```
In [59]: df4.reindex(['a', 'b', 'c', 'd'])
```

```
Out[59]:
```

	Ohio	Texas	California
a	0.0	1.0	2.0
b	NaN	NaN	NaN
c	3.0	4.0	5.0
d	6.0	7.0	8.0

```
In [60]: # for reindexing columns

df4.reindex(columns = ['Utah', 'Ohio', 'Texas'])
```

```
Out[60]:
```

	Utah	Ohio	Texas
a	NaN	0	1
c	NaN	3	4
d	NaN	6	7

## Dropping entries from an Axis

```
In [61]: # For Series
```

```
df1 = pd.Series(np.arange(5), index=['a', 'b', 'c', 'd', 'e'])  
df1
```

```
Out[61]: a    0  
        b    1  
        c    2  
        d    3  
        e    4  
        dtype: int64
```

```
In [62]: df1.drop('b')
```

```
Out[62]: a    0  
        c    2  
        d    3  
        e    4  
        dtype: int64
```

```
In [63]: df1.drop(['a', 'c'])
```

```
Out[63]: b    1  
        d    3  
        e    4  
        dtype: int64
```

```
In [64]: # For Data Frame
```

```
df2 = pd.DataFrame(np.arange(16).reshape((4, 4)),  
                    index=['Ohio', 'Colorado', 'Utah', 'New York'],  
                    columns=['one', 'two', 'three', 'four'])  
df2
```

```
Out[64]:
```

	one	two	three	four
Ohio	0	1	2	3
Colorado	4	5	6	7
Utah	8	9	10	11
New York	12	13	14	15

```
In [65]: # Default axis is rows (0)  
df2.drop('Ohio')
```

```
Out[65]:
```

	one	two	three	four
Colorado	4	5	6	7
Utah	8	9	10	11
New York	12	13	14	15

```
In [66]: df2.drop(['Colorado', 'Ohio'])
```

```
Out[66]:
```

	one	two	three	four
Utah	8	9	10	11
New York	12	13	14	15

```
In [67]: # From dropping columns

df2.drop('two', axis='columns')
```

```
Out[67]:
```

	one	three	four
Ohio	0	2	3
Colorado	4	6	7
Utah	8	10	11
New York	12	14	15

```
In [68]: df2.drop(['two', 'four'], axis=1)
```

```
Out[68]:
```

	one	three
Ohio	0	2
Colorado	4	6
Utah	8	10
New York	12	14

```
In [69]: df2.drop(['two', 'four'], axis=1, inplace = True)
df2
```

```
Out[69]:
```

	one	three
Ohio	0	2
Colorado	4	6
Utah	8	10
New York	12	14

## Indexing, Selection, and Filtering

```
In [70]: df1 = pd.Series(np.arange(10,14), index=['a', 'b', 'c', 'd'])
df1
```

```
Out[70]: a    10
b     11
c     12
d     13
dtype: int64
```

```
In [71]: df1['c']
```

```
Out[71]: 12
```

```
In [72]: df1[2]
```

```
Out[72]: 12
```

```
In [73]: df1[1:3]
```

```
Out[73]: b    11  
         c    12  
         dtype: int64
```

```
In [74]: # inclusive end-point
```

```
df1['b':'d']
```

```
Out[74]: b    11  
         c    12  
         d    13  
         dtype: int64
```

```
In [75]: df1[[3,1]]
```

```
Out[75]: d    13  
         b    11  
         dtype: int64
```

```
In [76]: df1[['d', 'b']]
```

```
Out[76]: d    13  
         b    11  
         dtype: int64
```

```
In [77]: df1[df1 < 12]
```

```
Out[77]: a    10  
         b    11  
         dtype: int64
```

```
In [78]: df1['b':'d'] = 50  
df1
```

```
Out[78]: a    10  
         b    50  
         c    50  
         d    50  
         dtype: int64
```

```
In [79]: # For Data Frame
```

```
df2 = pd.DataFrame(np.arange(16).reshape((4, 4)),  
                    index=['Ohio', 'Colorado', 'Utah', 'New York'],  
                    columns=['one', 'two', 'three', 'four'])  
df2
```

```
Out[79]:
```

	one	two	three	four
Ohio	0	1	2	3
Colorado	4	5	6	7
Utah	8	9	10	11
New York	12	13	14	15

```
In [80]: df2['two']
```

```
Out[80]: Ohio      1  
Colorado    5  
Utah        9  
New York   13  
Name: two, dtype: int64
```

```
In [81]: df2[['two', 'one']]
```

```
Out[81]:
```

	two	one
Ohio	1	0
Colorado	5	4
Utah	9	8
New York	13	12

```
In [82]: # Special cases
```

```
df2[:2]
```

```
Out[82]:
```

	one	two	three	four
Ohio	0	1	2	3
Colorado	4	5	6	7

```
In [83]: df2['three'] < 10
```

```
Out[83]: Ohio      True  
Colorado    True  
Utah      False  
New York   False  
Name: three, dtype: bool
```

```
In [84]: df2[df2['three'] < 10]
```

```
Out[84]:
```

	one	two	three	four
Ohio	0	1	2	3
Colorado	4	5	6	7

```
In [85]: df2
```

```
Out[85]:
```

	one	two	three	four
Ohio	0	1	2	3
Colorado	4	5	6	7
Utah	8	9	10	11
New York	12	13	14	15

```
In [86]: df2[df2 < 10] = -1
df2
```

Out[86]:

	one	two	three	four
Ohio	-1	-1	-1	-1
Colorado	-1	-1	-1	-1
Utah	-1	-1	10	11
New York	12	13	14	15

### Selecting with loc and iloc

- for DataFrame label-indexing on the rows
- loc (using axis labels)
- iloc (using integer index)

```
In [87]: df2 = pd.DataFrame(np.arange(16).reshape((4, 4)),
                           index=['Ohio', 'Colorado', 'Utah', 'New York'],
                           columns=['one', 'two', 'three', 'four'])
df2
```

Out[87]:

	one	two	three	four
Ohio	0	1	2	3
Colorado	4	5	6	7
Utah	8	9	10	11
New York	12	13	14	15

```
In [88]: df2.loc['Colorado']
```

Out[88]: one 4  
two 5  
three 6  
four 7  
Name: Colorado, dtype: int64

```
In [89]: df2.loc['Colorado', ['two', 'four']]
```

Out[89]: two 5  
four 7  
Name: Colorado, dtype: int64

```
In [90]: df2.iloc[1]
```

Out[90]: one 4  
two 5  
three 6  
four 7  
Name: Colorado, dtype: int64

```
In [91]: df2.iloc[1, [1, 3]]
```

Out[91]: two 5  
four 7  
Name: Colorado, dtype: int64



```
In [92]: df2.iloc[[1, 2]]
```

```
Out[92]:
```

	one	two	three	four
Colorado	4	5	6	7
Utah	8	9	10	11

```
In [93]: df2.iloc[[1, 2], [1, 3]]
```

```
Out[93]:
```

	two	four
Colorado	5	7
Utah	9	11

```
In [94]: df2.loc[: 'Utah']
```

```
Out[94]:
```

	one	two	three	four
Ohio	0	1	2	3
Colorado	4	5	6	7
Utah	8	9	10	11

```
In [95]: df2.loc[: 'Utah', ['two', 'three']]
```

```
Out[95]:
```

	two	three
Ohio	1	2
Colorado	5	6
Utah	9	10

```
In [96]: df2.iloc[:, :3]
```

```
Out[96]:
```

	one	two	three
Ohio	0	1	2
Colorado	4	5	6
Utah	8	9	10
New York	12	13	14

```
In [97]: df2.iloc[:, :3][df2.three > 5]
```

```
Out[97]:
```

	one	two	three
Colorado	4	5	6
Utah	8	9	10
New York	12	13	14

## Function application and mapping

```
In [98]: df1 = pd.DataFrame(np.random.randn(4, 3), columns=list('abc'),
                           index=['Utah', 'Ohio', 'Texas', 'Oregon'])
df1
```

Out[98]:

	a	b	c
Utah	-0.578600	1.651437	-2.426679
Ohio	-0.428913	1.265936	-0.866740
Texas	-0.678886	-0.094709	1.491390
Oregon	-0.638902	-0.443982	-0.434351

```
In [99]: np.abs(df1)
```

Out[99]:

	a	b	c
Utah	0.578600	1.651437	2.426679
Ohio	0.428913	1.265936	0.866740
Texas	0.678886	0.094709	1.491390
Oregon	0.638902	0.443982	0.434351

```
In [100]: # apply a function on 1-D arrays to each column or row
```

```
In [101]: # default axis = 'rows'

df1.apply(lambda x: x.max() - x.min())
```

```
Out[101]: a    0.249974
          b    2.095418
          c    3.918069
          dtype: float64
```

```
In [102]: # invoke once per row

df1.apply(lambda x: x.max() - x.min(), axis = 'columns')
```

```
Out[102]: Utah    4.078116
          Ohio    2.132677
          Texas    2.170276
          Oregon    0.204551
          dtype: float64
```

```
In [103]: df1
```

Out[103]:

	a	b	c
Utah	-0.578600	1.651437	-2.426679
Ohio	-0.428913	1.265936	-0.866740
Texas	-0.678886	-0.094709	1.491390
Oregon	-0.638902	-0.443982	-0.434351

```
In [104]: # function returning multiple values
```

```
df1.apply(lambda x: pd.Series([x.min(), x.max()], index = ['min', 'max']))
```

```
Out[104]:
```

	a	b	c
min	-0.678886	-0.443982	-2.426679
max	-0.428913	1.651437	1.491390

```
In [105]: df1.apply(lambda x: pd.Series([x.min(), x.max()], index = ['min', 'max']),  
axis='columns')
```

```
Out[105]:
```

	min	max
Utah	-2.426679	1.651437
Ohio	-0.866740	1.265936
Texas	-0.678886	1.491390
Oregon	-0.638902	-0.434351

## Sorting

- sort lexicographically by row or column index

```
In [106]: # Series
```

```
df1 = pd.Series(np.arange(10,14), index=['d', 'a', 'b', 'c'])  
df1
```

```
Out[106]: d    10  
a     11  
b     12  
c     13  
dtype: int64
```

```
In [107]: df2 = df1.sort_index()  
df2
```

```
Out[107]: a     11  
b     12  
c     13  
d     10  
dtype: int64
```

```
In [108]: df2.sort_values()
```

```
Out[108]: d     10  
a     11  
b     12  
c     13  
dtype: int64
```

```
In [109]: # DataFrame

df1 = pd.DataFrame(np.arange(8).reshape((2, 4)),
                    index=['three', 'one'],
                    columns=['d', 'a', 'b', 'c'])

df1
```

```
Out[109]:
```

	d	a	b	c
three	0	1	2	3
one	4	5	6	7

```
In [110]: df1.sort_index()
```

```
Out[110]:
```

	d	a	b	c
one	4	5	6	7
three	0	1	2	3

```
In [111]: df1.sort_index(axis=1)
```

```
Out[111]:
```

	a	b	c	d
three	1	2	3	0
one	5	6	7	4

```
In [112]: df1
```

```
Out[112]:
```

	d	a	b	c
three	0	1	2	3
one	4	5	6	7

```
In [113]: df1.sort_values(by='b', ascending=False)
```

```
Out[113]:
```

	d	a	b	c
one	4	5	6	7
three	0	1	2	3

```
In [114]: df1.sort_values(by='one', axis=1, ascending=False)
```

```
Out[114]:
```

	c	b	a	d
three	3	2	1	0
one	7	6	5	4

## Axis indices with duplicate labels

```
In [115]: # Series
```

```
df1 = pd.Series(np.arange(10,15), index=['a', 'a', 'b', 'b', 'c'])  
df1
```

```
Out[115]: a    10  
          a    11  
          b    12  
          b    13  
          c    14  
          dtype: int64
```

```
In [116]: df1['b']
```

```
Out[116]: b    12  
          b    13  
          dtype: int64
```

```
In [117]: df1.index.is_unique
```

```
Out[117]: False
```

```
In [118]: # DataFrame
```

```
df2 = pd.DataFrame(np.random.randint(60, 90, (4, 3)), index=['a', 'a', 'b', 'b'])  
df2
```

```
Out[118]:
```

	0	1	2
a	67	62	80
a	75	84	89
b	76	67	69
b	63	88	88

```
In [119]: df2.loc['b']
```

```
Out[119]:
```

	0	1	2
b	76	67	69
b	63	88	88

## Descriptive Statistics

```
In [120]: df1 = pd.DataFrame([[1.5, np.nan], [7.5, -5.5],  
                               [np.nan, np.nan], [1.0, -4.5]],  
                               index=['a', 'b', 'c', 'd'],  
                               columns=['one', 'two'])  
df1
```

```
Out[120]:
```

	one	two
a	1.5	NaN
b	7.5	-5.5
c	NaN	NaN
d	1.0	-4.5

```
In [121]: df1.sum()
```

```
Out[121]: one      10.0  
two      -10.0  
dtype: float64
```

```
In [122]: df1.sum(axis=0)
```

```
Out[122]: one      10.0  
two      -10.0  
dtype: float64
```

```
In [123]: df1.sum(axis='rows')
```

```
Out[123]: one      10.0  
two      -10.0  
dtype: float64
```

```
In [124]: df1.sum(axis=1)
```

```
Out[124]: a      1.5  
b      2.0  
c      0.0  
d     -3.5  
dtype: float64
```

```
In [125]: df1.sum(axis='columns')
```

```
Out[125]: a      1.5  
b      2.0  
c      0.0  
d     -3.5  
dtype: float64
```

```
In [126]: df1
```

```
Out[126]:
```

	one	two
a	1.5	NaN
b	7.5	-5.5
c	NaN	NaN
d	1.0	-4.5

## idxmax, idxmin

- index labels of maximum and minimum values

## argmax, argmin (Series)

- index locations of maximum and minimum values for a Series

```
In [127]: print(df1)

df1.idxmax()
```

```
   one  two
a  1.5  NaN
b  7.5 -5.5
c  NaN  NaN
d  1.0 -4.5
```

```
Out[127]: one    b
         two    d
         dtype: object
```

```
In [128]: df1.idxmax(axis='columns')
```

```
Out[128]: a    one
         b    one
         c    NaN
         d    one
         dtype: object
```

## accumulations

- cumsum, cumprod, cummin, cummax

```
In [129]: print(df1)

df1.cumsum()
```

```
   one  two
a  1.5  NaN
b  7.5 -5.5
c  NaN  NaN
d  1.0 -4.5
```

```
Out[129]:
```

	one	two
<b>a</b>	1.5	NaN
<b>b</b>	9.0	-5.5
<b>c</b>	NaN	NaN
<b>d</b>	10.0	-10.0

```
In [130]: df1.describe()
```

Out[130]:

	one	two
count	3.000000	2.000000
mean	3.333333	-5.000000
std	3.617089	0.707107
min	1.000000	-5.500000
25%	1.250000	-5.250000
50%	1.500000	-5.000000
75%	4.500000	-4.750000
max	7.500000	-4.500000

```
In [131]: df1
```

Out[131]:

	one	two
a	1.5	NaN
b	7.5	-5.5
c	NaN	NaN
d	1.0	-4.5

```
In [132]: np.random.seed(123)
df2 = pd.DataFrame(np.random.randint(60, 90, (4, 3)),
                    index=['a', 'b', 'c', 'd'],
                    columns = ['one', 'two', 'three'])
df2
```

Out[132]:

	one	two	three
a	73	62	88
b	62	66	77
c	79	70	87
d	85	82	61

```
In [133]: df2.diff()
```

Out[133]:

	one	two	three
a	NaN	NaN	NaN
b	-11.0	4.0	-11.0
c	17.0	4.0	10.0
d	6.0	12.0	-26.0



```
In [134]: df2.diff(axis='columns')
```

Out[134]:

	one	two	three
a	NaN	-11.0	26.0
b	NaN	4.0	11.0
c	NaN	-9.0	17.0
d	NaN	-3.0	-21.0

```
In [135]: df2
```

Out[135]:

	one	two	three
a	73	62	88
b	62	66	77
c	79	70	87
d	85	82	61

```
In [136]: df2.pct_change()
```

Out[136]:

	one	two	three
a	NaN	NaN	NaN
b	-0.150685	0.064516	-0.125000
c	0.274194	0.060606	0.129870
d	0.075949	0.171429	-0.298851

```
In [137]: df2['one'].cov(df2['two'])
```

Out[137]: 62.666666666666664

```
In [138]: df2['one'].corr(df2['two'])
```

Out[138]: 0.7392185280134137

```
In [139]: df2.cov()
```

Out[139]:

	one	two	three
one	96.250000	62.666667	-46.916667
two	62.666667	74.666667	-93.333333
three	-46.916667	-93.333333	156.916667


```
In [140]: df2.corr()
```

Out[140]:


	one	two	three
one	1.000000	0.739219	-0.381762
two	0.739219	1.000000	-0.862261
three	-0.381762	-0.862261	1.000000

```
In [141]: import matplotlib.pyplot as plt
```

```
In [142]: plt.scatter(df2['one'], df2['two']);
```



```
In [143]: plt.scatter(df2['two'], df2['three']);
```



## Unique values and value counts

```
In [144]: np.random.seed(123)
scores = np.random.randint(60, 70, 10)

a = pd.Series(scores)
a
```

```
Out[144]: 0    62
          1    62
          2    66
          3    61
          4    63
          5    69
          6    66
          7    61
          8    60
          9    61
          dtype: int64
```

```
In [145]: a.unique()
```

```
Out[145]: array([62, 66, 61, 63, 69, 60])
```

```
In [146]: a.value_counts()
```

```
Out[146]: 61    3
          62    2
          66    2
          63    1
          60    1
          69    1
          dtype: int64
```

```
In [147]: a.values
```

```
Out[147]: array([62, 62, 66, 61, 63, 69, 66, 61, 60, 61])
```

```
In [148]: pd.value_counts(a)
```

```
Out[148]: 61    3
          62    2
          66    2
          63    1
          60    1
          69    1
          dtype: int64
```

```
In [149]: pd.value_counts(a.values)
```

```
Out[149]: 61    3
          62    2
          66    2
          63    1
          60    1
          69    1
          dtype: int64
```

```
In [150]: pd.value_counts(a.values, sort=False)
```

```
Out[150]: 66    2
          69    1
          60    1
          61    3
          62    2
          63    1
          dtype: int64
```

```
In [151]: a.unique()
```

```
Out[151]: array([62, 66, 61, 63, 69, 60])
```

```
In [152]: pd.Index(a.unique()).get_indexer(a)
```

```
Out[152]: array([0, 0, 1, 2, 3, 4, 1, 2, 5, 2])
```

```
In [153]: np.random.seed(321)
df2 = pd.DataFrame(np.random.randint(60, 70, (10, 4)),
                    columns = ['Q1', 'Q2', 'Q3', 'Q4'])
df2
```

```
Out[153]:
```

	Q1	Q2	Q3	Q4
0	64	69	68	61
1	68	68	64	65
2	68	63	65	61
3	64	66	65	67
4	67	62	62	63
5	69	62	61	62
6	61	61	60	64
7	64	63	60	63
8	67	64	65	67
9	60	68	67	61

```
In [154]: df2.apply(pd.value_counts)
```

```
Out[154]:
```

	Q1	Q2	Q3	Q4
60	1.0	NaN	2.0	NaN
61	1.0	1.0	1.0	3.0
62	NaN	2.0	1.0	1.0
63	NaN	2.0	NaN	2.0
64	3.0	1.0	1.0	1.0
65	NaN	NaN	3.0	1.0
66	NaN	1.0	NaN	NaN
67	2.0	NaN	1.0	2.0
68	2.0	2.0	1.0	NaN
69	1.0	1.0	NaN	NaN

```
In [155]: df2.apply(pd.value_counts).dropna()
```

```
Out[155]:
```

	Q1	Q2	Q3	Q4
61	1.0	1.0	1.0	3.0
64	3.0	1.0	1.0	1.0

```
In [156]: df2.apply(pd.value_counts).fillna(0)
```

Out[156]:

	Q1	Q2	Q3	Q4
<b>60</b>	1.0	0.0	2.0	0.0
<b>61</b>	1.0	1.0	1.0	3.0
<b>62</b>	0.0	2.0	1.0	1.0
<b>63</b>	0.0	2.0	0.0	2.0
<b>64</b>	3.0	1.0	1.0	1.0
<b>65</b>	0.0	0.0	3.0	1.0
<b>66</b>	0.0	1.0	0.0	0.0
<b>67</b>	2.0	0.0	1.0	2.0
<b>68</b>	2.0	2.0	1.0	0.0
<b>69</b>	1.0	1.0	0.0	0.0

In [ ]: