

Classification - Support Vector Machines (SVM)

- Support vector machines (SVMs) are a set of supervised learning methods used for classification, regression and outliers detection.

```
In [1]: import numpy as np
import matplotlib.pyplot as plt

# use seaborn plotting defaults
import seaborn as sns
sns.set_style('whitegrid')
```

```
In [2]: from sklearn.datasets.samples_generator import make_blobs
```

```
In [3]: # default n_features = 2

X, y = make_blobs(n_samples=50, centers=2,
                  random_state=0, cluster_std=0.60)
```


```
In [4]: x[:0.5, 1]
```

```
Out[4]: array([[1.41281595, 1.5303347 ],
               [1.81336135, 1.6311307 ],
               [1.43289271, 4.37679234],
               [1.87271752, 4.18069237],
               [2.09517785, 1.0791468 ]])
```

```
In [5]: v[:0.5, 1]
```

```
Out[5]: array([1, 1, 0, 0, 1])
```

```
In [6]: plt.scatter(X[:, 0], X[:, 1], c=y, s=50, cmap='RdBu').
```




```
In [7]: # Draw a line separating the data - more than one possible lines
# Sample lines (y = mx + b)

xfit = np.linspace(-1, 3.5)
plt.scatter(X[:, 0], X[:, 1], c=y, s=50, cmap='RdBu')
plt.plot([0.6], [2.1], 'x', color='red', markeredgewidth=2, markersize=10)

for m, b in [(1, 0.65), (0.5, 1.6), (-0.2, 2.9)]:
    plt.plot(xfit, m * xfit + b, '-k')

plt.xlim(-1, 3.5):
```



Depending on which line you choose, a new data point (e.g., the one marked by the "X" in this plot) will be assigned a different label


Support Vector Machines - Maximizing the Margin

Rather than simply drawing a zero-width line between the classes, we can draw around each line a margin of some width, up to the nearest point.

```
In [8]: xfit = np.linspace(-1, 3.5)
plt.scatter(X[:, 0], X[:, 1], c=y, s=50, cmap='RdBu')

for m, b, d in [(1, 0.65, 0.33), (0.5, 1.6, 0.55), (-0.2, 2.9, 0.2)]:
    yfit = m * xfit + b
    plt.plot(xfit, yfit, '-k')
    plt.fill_between(xfit, yfit - d, yfit + d, edgecolor='none',
                    color='g', alpha=0.2)

plt.xlim(-1, 3.5).
```



In support vector machines, the line that maximizes this margin is the one we will choose as the optimal model.

Support vector machines are an example of such a maximum margin estimator.

```
In [9]: from sklearn.svm import SVC      # "Support vector classifier"
```

```
In [10]: model = SVC(kernel='linear')
```

```
model.fit(X, y)
```

```
Out[10]: SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
            decision_function_shape='ovr', degree=3, gamma='auto_deprecated',
            kernel='linear', max_iter=-1, probability=False, random_state=None,
            shrinking=True, tol=0.001, verbose=False)
```

```
In [11]: model.support_vectors
```

```
Out[11]: array([[0.44359863, 3.11530945],
                [2.33812285, 3.43116792],
                [2.06156753, 1.96918596]])
```

```
In [12]: def plot_svc_decision_function(model, ax=None, plot_support=True):
```

```
    """Plot the decision function for a 2D SVC"""
```

```
    if ax is None:
```

```
        ax = plt.gca()
```

```
    xlim = ax.get_xlim()
```

```
    ylim = ax.get_ylim()
```

```
    # create grid to evaluate model
```

```
    x = np.linspace(xlim[0], xlim[1], 30)
```

```
    y = np.linspace(ylim[0], ylim[1], 30)
```

```
    Y, X = np.meshgrid(y, x)
```

```
    xy = np.vstack([X.ravel(), Y.ravel()]).T
```

```
    P = model.decision_function(xy).reshape(X.shape)
```

```
    # plot decision boundary and margins
```

```
    ax.contour(X, Y, P, colors='k',
```

```
               levels=[-1, 0, 1],
```

```
               alpha=0.5,
```

```
               linestyles=['--', '-', '--'])
```

```
    # plot support vectors
```

```
    if plot_support:
```

```
        ax.scatter(model.support_vectors[:, 0],
```


```
                  model.support_vectors[:, 1],
```

```
                  s=100, linewidth=1, facecolors='g');
```

```
    ax.set_xlim(xlim)
```

```
    ax.set_ylim(ylim)
```

```
In [13]: plt.scatter(X[:, 0], X[:, 1], c=y, s=50, cmap='RdBu')
plt_svc.decision_function(model):
```




```
In [ ]:
```

```
In [14]: from sklearn.datasets.samples_generator import make_circles
```

```
In [15]: X, y = make_circles(100, factor=.1, noise=.1)


plt.scatter(X[:, 0], X[:, 1], c=y, s=50, cmap='RdBu')
plt.xlim(-1.5, 1.5)
plt.ylim(-1.5, 1.5).
```



In [16]: `# Linear kernel will not work`

```
model2 = SVC(kernel='linear').fit(X, y)
```

In [17]: `plt.scatter(X[:, 0], X[:, 1], c=y, s=50, cmap='RdBu')
plt.xlim(-1.5, 1.5)
plt.ylim(-1.5, 1.5)
plt.plot_svm_decision_function(model2, plot_support=False)`




In Scikit-Learn, we can apply kernelized SVM simply by changing our linear kernel to an RBF (radial basis function) kernel, using the kernel model hyperparameter:

In [18]: `model3 = SVC(kernel='rbf', C=1E6, gamma='auto')
model3.fit(X, y)`

Out[18]: `SVC(C=1000000.0, cache_size=200, class_weight=None, coef0=0.0,
decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
max_iter=-1, probability=False, random_state=None, shrinking=True,
tol=0.001, verbose=False)`

```
In [19]: plt.scatter(X[:, 0], X[:, 1], c=y, s=50, cmap='RdBu')
plt.xlim(-1.5, 1.5)
plt.ylim(-1.5, 1.5)
plt.show()
```



```
In [20]: model3.support_vectors_
```

```
Out[20]: array([[ 0.87458364, -0.32529304],
 [-0.88285475,  0.26816679],
 [ 0.0503305 ,  0.84845701],
 [ 0.00243228, -0.85618398],
 [ 0.59933641, -0.6709688 ],
 [ 0.5734485 ,  0.65138603],
 [-0.71816916, -0.40812948],
 [-0.22407102, -0.21182966]])
```

```
In [ ]:
```

Case Study - Face Recognition

```
In [21]: from sklearn.datasets import fetch_lfw_people
```

```
In [22]: faces = fetch_lfw_people(min_faces_per_person=70)
```

```
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/externals/joblib/__init___.py:15: DeprecationWarning: sklearn.externals.joblib is deprecated in 0.21 and will be removed in 0.23. Please import this functionality directly from joblib, which can be installed with: pip install joblib. If this warning is raised when loading pickled models, you may need to re-serialize those models with scikit-learn 0.21+.
  warnings.warn(msg, category=DeprecationWarning)
```

```
In [23]: print(faces.target_names)
         print(faces.images.shape)
```

```
['Ariel Sharon' 'Colin Powell' 'Donald Rumsfeld' 'George W Bush'
 'Gerhard Schroeder' 'Hugo Chavez' 'Tony Blair']
(1288, 62, 47)
```

```
In [24]: faces.images[0].shape
```

```
Out[24]: (62, 47)
```



```
In [25]: # Plot a few of the faces

fig, ax = plt.subplots(3, 5, figsize=(12,8))
for i, axi in enumerate(ax.flat):
    axi.imshow(faces.images[i], cmap='bone')
    axi.set(xticks=[], yticks=[],
            xlabel=faces.target_names[faces.target[i]])
```



Each image contains $[62 \times 47]$ or nearly 3,000 pixels. Not effective to use each pixel value as a feature. We will use a principal component analysis to extract 150 fundamental components to feed into our support vector machine classifier

```
In [26]: from sklearn.svm import SVC
from sklearn.decomposition import PCA
from sklearn.pipeline import Pipeline
```

```
In [27]: pca = PCA(n_components=150, svd_solver='randomized', whiten=True)

svc = SVC(kernel='rbf', gamma=0.005, C=1000, class_weight='balanced')

model = Pipeline([('pca', pca),
                  ('svc', svc)])
```

For testing the classifier output, we will split the data into a training and testing set

```
In [28]: from sklearn.model_selection import train_test_split
```

```
In [29]: Xtrain, Xtest, ytrain, ytest = train_test_split(faces.data, faces.target,
                                                    random_state=42) # default test size 0.25
```

```
In [30]: np.unique(faces.target)
```

```
Out[30]: array([0, 1, 2, 3, 4, 5, 6])
```

```
In [31]: model.fit(Xtrain, ytrain)
```

```
Out[31]: Pipeline(memory=None,
                  steps=[('pca',
                        PCA(copy=True, iterated_power='auto', n_components=150,
                            random_state=None, svd_solver='randomized', tol=0.0,
                            whiten=True)),
                        ('svc',
                        SVC(C=1000, cache_size=200, class_weight='balanced', coef0=0.0,
                            decision_function_shape='ovr', degree=3, gamma=0.005,
                            kernel='rbf', max_iter=-1, probability=False,
                            random_state=None, shrinking=True, tol=0.001,
                            verbose=False))],
                  verbose=False)
```

```
In [32]: yfit = model.predict(Xtest)
```

```
In [33]: from sklearn.metrics import confusion_matrix
```

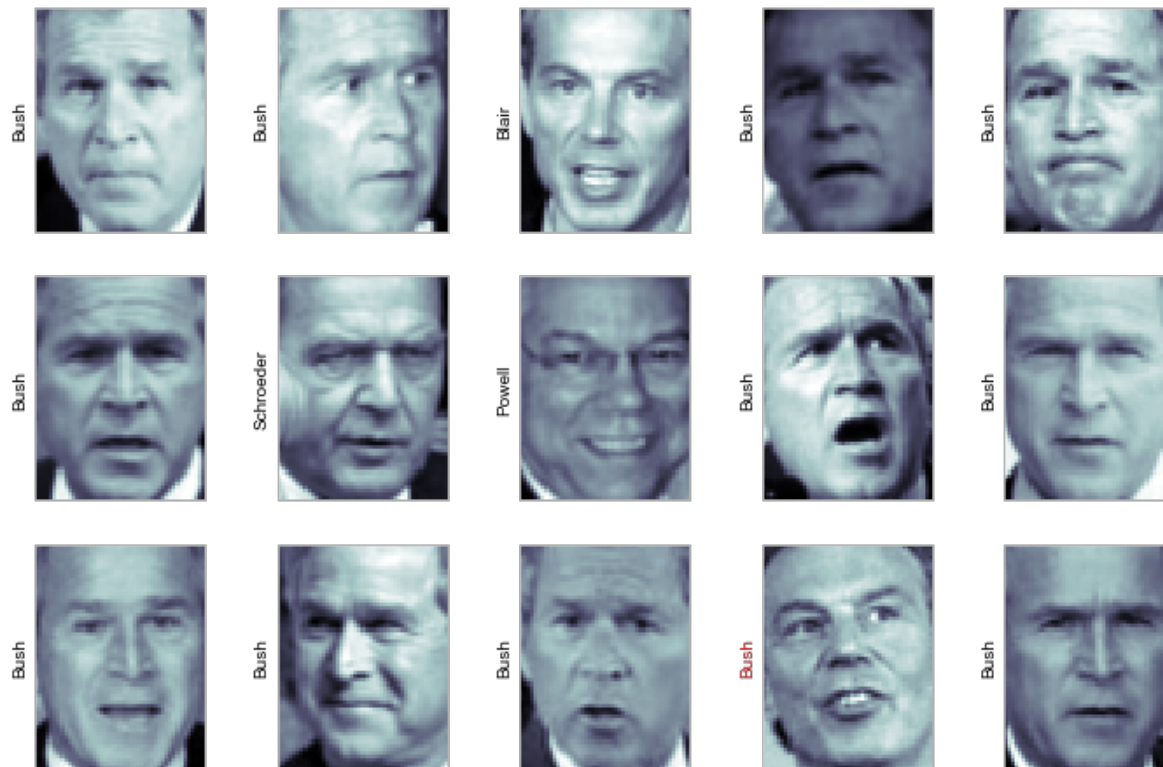
```
In [34]: confusion_matrix(ytest, yfit)
```

```
Out[34]: array([[ 6,  2,  0,  5,  0,  0,  0],
 [ 1, 52,  0,  7,  0,  0,  0],
 [ 0,  2, 17,  8,  0,  0,  0],
 [ 0,  3,  0, 143,  0,  0,  0],
 [ 0,  1,  0,  4, 20,  0,  0],
 [ 0,  3,  0,  3,  1,  8,  0],
 [ 1,  1,  2,  4,  0,  0, 28]])
```


```
In [35]: fig, ax = plt.subplots(3, 5, figsize=(12,8))

for i, axi in enumerate(ax.flat):
    axi.imshow(Xtest[i].reshape(62,47), cmap='bone')
    axi.set(xticks=[], yticks=[])
    axi.set_ylabel(faces.target_names[yfit[i]].split()[-1],
                  color='black' if yfit[i] == ytest[i] else 'red')
fig.suptitle('Predicted Names; Incorrect Labels in Red' size=14).
```

Predicted Names; Incorrect Labels in Red



```
In [36]: mat = confusion_matrix(ytest, yfit)
sns.heatmap(mat.T, square=True, annot=True, fmt='d', cbar=False,
            xticklabels=faces.target_names,
            yticklabels=faces.target_names)
plt.xlabel('true label')
plt.ylabel('predicted label').
```



GridSearchCV

```
In [37]: from sklearn.model_selection import GridSearchCV
```

```
In [38]: Xtrain_pca = pca.transform(Xtrain)
```

```
In [39]: param_grid = {'C': [1e3, 5e3, 1e4, 5e4, 1e5],
                      'gamma': [0.0001, 0.0005, 0.001, 0.005, 0.01, 0.1], }

clf = GridSearchCV(SVC(kernel='rbf', class_weight='balanced'),
                  param_grid, cv=5) # 5 Fold cross validation
```

```
In [40]: clf = clf.fit(Xtrain_pca, ytrain)
```

```
In [41]: print("Best estimator found by grid search:")
print(cgf.best_estimator_)
Best estimator found by grid search:
SVC(C=1000.0, cache_size=200, class_weight='balanced', coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma=0.005, kernel='rbf',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)
```

PCA - # of features

```
In [42]: pca = PCA(n_components=10, svd_solver='randomized', whiten=True)
pca.fit(faces.data)
```

```
Out[42]: PCA(copy=True, iterated_power='auto', n_components=10, random_state=None,
    svd_solver='randomized', tol=0.0, whiten=True)
```

```
In [43]: components = pca.transform(faces.data)
projected = pca.inverse_transform(components)
```

```
In [44]: # Plot the results

fig, ax = plt.subplots(2, 10, figsize=(10, 2.5),
    subplot_kw={'xticks':[], 'yticks':[]},
    gridspec_kw=dict(hspace=0.1, wspace=0.1))

for i in range(10):
    ax[0, i].imshow(faces.data[i].reshape(62, 47), cmap='bone')
    ax[1, i].imshow(projected[i].reshape(62, 47), cmap='bone')

ax[0, 0].set_ylabel('full-dim\ninput')
ax[1, 0].set_ylabel('10-dim\nreconstruction')
```




```
In [45]: pca.explained_variance_ratio_
```

```
Out[45]: array([0.1841573 , 0.1474062 , 0.07128289, 0.05923904, 0.05049089,
                0.03009569, 0.02451012, 0.02092452, 0.02018103, 0.01896057],
                dtype=float32)
```

```
In [46]: np.cumsum(pca.explained_variance_ratio_)
```

```
Out[46]: array([0.1841573 , 0.3315635 , 0.4028464 , 0.46208543, 0.51257634,
                0.54267204, 0.5671822 , 0.5881067 , 0.6082877 , 0.6272483 ],
                dtype=float32)
```


```
In [47]: plt.plot(np.cumsum(pca.explained_variance_ratio_))
plt.xlabel('number of components')
plt.ylabel('cumulative explained variance').
```



```
In [48]: pca = PCA(n_components=250, svd_solver='randomized', whiten=True)
pca.fit(faces_data)
```

```
Out[48]: PCA(copy=True, iterated_power='auto', n_components=250, random_state=None,
            svd_solver='randomized', tol=0.0, whiten=True)
```

```
In [49]: plt.plot(np.cumsum(pca.explained_variance_ratio_))
plt.xlabel('number of components')
plt.ylabel('cumulative explained variance').
```



```
In [50]: components = pca.transform(faces.data)
projected = pca.inverse_transform(components)
```

```
In [51]: # Plot the results

fig, ax = plt.subplots(2, 10, figsize=(10, 2.5),
                        subplot_kw={'xticks':[], 'yticks':[]},
                        gridspec_kw=dict(hspace=0.1, wspace=0.1))

for i in range(10):
    ax[0, i].imshow(faces.data[i].reshape(62, 47), cmap='bone')
    ax[1, i].imshow(projected[i].reshape(62, 47), cmap='bone')

ax[0, 0].set_ylabel('full-dim\ninput')
ax[1, 0].set_ylabel('250-dim\nreconstruction').
```



In []: