

Language Classification

```
In [1]: from sklearn.feature_extraction.text import TfidfVectorizer

        from sklearn.svm import LinearSVC

        from sklearn.pipeline import Pipeline
        from sklearn.datasets import load_files
        from sklearn.model_selection import train_test_split
        from sklearn import metrics
```

```
In [2]: # The training data folder

        languages_data_folder = 'data/languages/paragraphs'

        dataset = load_files(languages_data_folder)
```

```
In [3]: dataset.target_names
```

```
Out[3]: ['ar', 'de', 'en', 'es', 'fr', 'it', 'ja', 'nl', 'pl', 'pt', 'ru']
```

```
In [4]: # Split the dataset in training and test set:

        docs_train, docs_test, y_train, y_test = train_test_split(
            dataset.data, dataset.target, test_size=0.5, random_state=0)
```

- Build a vectorizer that splits strings into sequence of 1 to 3 characters instead of word tokens

```
In [5]: vectorizer = TfidfVectorizer(ngram_range=(1, 3), analyzer='char',
                                     use_idf=False)
```

- Build a vectorizer / classifier pipeline using the previous analyzer

```
In [6]: clf = Pipeline([
        ('vec', vectorizer),
        ('clf', LinearSVC(C=10000)),
    ])
```

- Fit the pipeline on the training set

```
In [7]: clf.fit(docs_train, y_train);
```

- Predict the outcome on the testing set

```
In [8]: y_predicted = clf.predict(docs_test)
```

```
In [9]: # Print the classification report

print(metrics.classification_report(
    y_test, y_predicted, target_names=dataset.target_names))
```

	precision	recall	f1-score	support
ar	1.00	1.00	1.00	11
de	1.00	1.00	1.00	82
en	1.00	1.00	1.00	68
es	1.00	1.00	1.00	66
fr	1.00	1.00	1.00	65
it	1.00	1.00	1.00	39
ja	1.00	1.00	1.00	35
nl	1.00	1.00	1.00	26
pl	1.00	1.00	1.00	21
pt	1.00	1.00	1.00	53
ru	1.00	1.00	1.00	29
accuracy			1.00	495
macro avg	1.00	1.00	1.00	495
weighted avg	1.00	1.00	1.00	495

In [10]: *# Show the confusion matrix*

```
cm = metrics.confusion_matrix(y_test, y_predicted)
print(cm)
```

```
[[11  0  0  0  0  0  0  0  0  0  0]
 [ 0 82  0  0  0  0  0  0  0  0  0]
 [ 0  0 68  0  0  0  0  0  0  0  0]
 [ 0  0  0 66  0  0  0  0  0  0  0]
 [ 0  0  0  0 65  0  0  0  0  0  0]
 [ 0  0  0  0  0 39  0  0  0  0  0]
 [ 0  0  0  0  0  0 35  0  0  0  0]
 [ 0  0  0  0  0  0  0 26  0  0  0]
 [ 0  0  0  0  0  0  0  0 21  0  0]
 [ 0  0  0  0  0  0  0  0  0 53  0]
 [ 0  0  0  0  0  0  0  0  0  0 29]]
```

In []:

In [11]: *# Predict the result on some short new sentences:*

```
sentences = [
    'This is a language detection test.',
    'Ceci est un test de d\xe9tection de la langue.',
    'Dies ist ein Test, um die Sprache zu erkennen.',
    'Questo \u00e8 un test di rilevamento della lingua.'
]

predicted = clf.predict(sentences)
```

In [12]: **for** s, p **in** zip(sentences, predicted):
 print("The language of {} is '{}'.format(s, dataset.target_names[p]))

```
The language of This is a language detection test. is 'en'
The language of Ceci est un test de d\u00e9tection de la langue. is 'es'
The language of Dies ist ein Test, um die Sprache zu erkennen. is 'de'
The language of Questo \u00e8 un test di rilevamento della lingua. is 'it'
```

In []: