Text mining



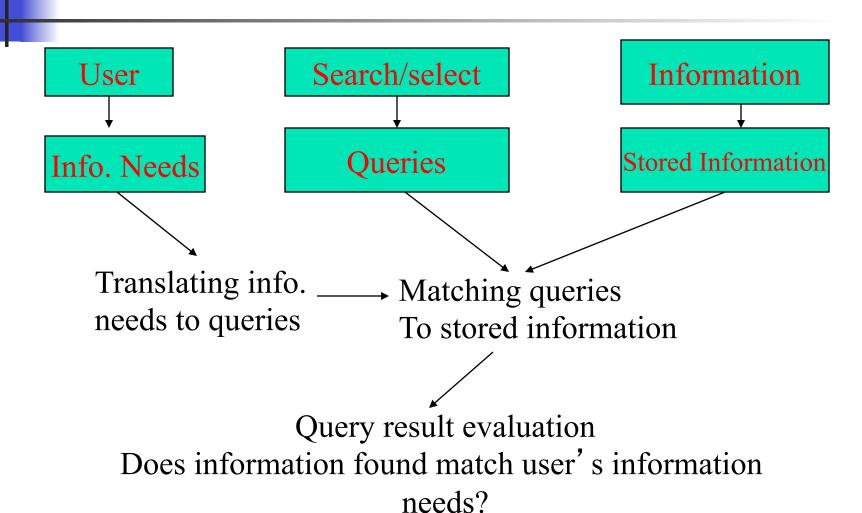


- It refers to data mining using text documents as data.
- There are many special techniques for pre-processing text documents to make them suitable for mining.
- Most of these techniques are from the field of "Information Retrieval".

Information Retrieval (IR)

- Conceptually, information retrieval (IR) is the study of finding needed information. i.e., IR helps users find information that matches their information needs.
- Historically, information retrieval is about document retrieval, emphasizing document as the basic unit.
- Technically, IR studies the acquisition, organization, storage, retrieval, and distribution of information.
- IR has become a center of focus in the Web era.

Information Retrieval





Text Processing

- Word (token) extraction
- Stop words
- Stemming
- Frequency counts

Stop words

- Many of the most frequently used words in English are worthless in IR and text mining – these words are called stop words.
 - the, of, and, to,
 - Typically about 400 to 500 such words
 - For an application, an additional domain specific stop words list may be constructed
 - Why do we need to remove stop words?
 - Reduce indexing (or data) file size
 - stopwords accounts 20-30% of total word counts.
 - Improve efficiency
 - stop words are not useful for searching or text mining
 - stop words always have a large number of hits

Stemming

Techniques used to find out the root/stem of a word:

E.g.,

user

users

used

using

engineering

engineered

engineer

stem: use engineer

Usefulness

- improving effectiveness of IR and text mining
 - matching similar words
- reducing indexing size
 - combing words with same roots may reduce indexing size as much as 40-50%.

Basic stemming methods

remove ending

- if a word ends with a consonant other than s, followed by an s, then delete s.
- if a word ends in es, drop the s.
- if a word ends in ing, delete the ing unless the remaining word consists only of one letter or of th.
- If a word ends with ed, preceded by a consonant, delete the ed unless this leaves only a single letter.
- **.....**

transform words

if a word ends with "ies" but not "eies" or "aies" then "ies --> y."

Frequency counts

- Counts the number of times a word occurred in a document.
- Counts the number of documents in a collection that contains a word.
- Using occurrence frequencies to indicate relative importance of a word in a document.
 - if a word appears often in a document, the document likely "deals with" subjects related to the word.

Vector Space Representation

- A document is represented as a vector:
 - (W₁, W₂, ..., W_n)
 - Binary:
 - Wi= 1 if the corresponding term i (often a word) is in the document
 - W_i= 0 if the term i is not in the document
 - TF: (Term Frequency)
 - Wi= tfi where tfi is the number of times the term occurred in the document
 - TF*IDF: (Inverse Document Frequency)
 - Wi =tfi*idfi=tfi*log(N/dfi)) where dfi is the number of documents contains term i, and N the total number of documents in the collection.

Vector Space and Document Similarity

- Each indexing term is a dimension. A indexing term is normally a word.
- Each document is a vector

$$D_i = (t_{i1}, t_{i2}, t_{i3}, t_{i4}, ... t_{in})$$

$$D_j = (t_{j1}, t_{j2}, t_{j3}, t_{j4}, ..., t_{jn})$$

Document similarity is defined as

Similarity (D_i, D_j) =
$$\frac{\sum_{k=1}^{n} t_{ik} * t_{jk}}{\sqrt{\sum_{k=1}^{n} t_{ik}^{2}} \times \sqrt{\sum_{k=1}^{n} t_{jk}^{2}}}$$



Query formats

- Query is a representation of the user's information needs
 - Normally a list of words.
- Query as a simple question in natural language
 - The system translates the question into executable queries
- Query as a document
 - "Find similar documents like this one"
 - The system defines what the similarity is

An Example

- A document Space is defined by three terms:
 - hardware, software, users
 - A set of documents are defined as:

```
- A1=(1, 0, 0), A2=(0, 1, 0), A3=(0, 0, 1)
```

•
$$A4=(1, 1, 0)$$
, $A5=(1, 0, 1)$, $A6=(0, 1, 1)$

•
$$A7=(1, 1, 1)$$
 $A8=(1, 0, 1)$. $A9=(0, 1, 1)$

- If the Query is "hardware and software"
- what documents should be retrieved?

An Example (cont.)

- In Boolean query matching:
 - document A4, A7 will be retrieved ("AND")
 - retrieved:A1, A2, A4, A5, A6, A7, A8, A9 ("OR")
- In similarity matching (cosine):
 - q=(1, 1, 0)
 - S(q, A1)=0.71, S(q, A2)=0.71, S(q, A3)=0
 - S(q, A4)=1, S(q, A5)=0.5, S(q, A6)=0.5
 - S(q, A7)=0.82, S(q, A8)=0.5, S(q, A9)=0.5
 - Document retrieved set (with ranking)=
 - {A4, A7, A1, A2, A5, A6, A8, A9}

Cosine Similarity

- Cosine Similarity is a technique that is derived from vector theory.
- In Information Retrieval, it is used to indicate (or measure)
 - the degree of similarity between two documents, or
 - between a document and a query.



- Keywords (Terms)
 - to describe the information content within a document.
- Vocabulary (Dictionary)
 - The total set of keywords
- Stopwords
 - Words which do not help to differentiate documents or which don't identify the information within a document
 - Discarded from the list of keywords

Examp

Car VanBus3 RoadHighway Bicycle0 CoachTrainStation Ticket

- Vector Inner Product (Dot Product)
 - is defined as the sum of the products of the vector components.

let
$$\mathbf{v} = \begin{pmatrix} 3 \\ -1 \\ 2 \end{pmatrix}$$
 and $\mathbf{w} = \begin{pmatrix} 1 \\ 2 \\ 1 \end{pmatrix}$

$$\mathbf{v}.\,\mathbf{w} = (3.1) + (-1.2) + (2.1) = 3$$

- Vector Length (Norm)
 - Inner product of vector with itself


$$||\mathbf{v}|| = (\mathbf{v}.\,\mathbf{v})^{1/2} \text{ or } \sqrt{(\mathbf{v}.\,\mathbf{v})}$$

$$\mathbf{v} = \begin{pmatrix} 3 \\ 4 \end{pmatrix}$$

$$\mathbf{v} \cdot \mathbf{v} = 3^2 + 4^2 = 25$$

- Take the square root
 - Length = 5

Cosine Similarity



$$\cos\theta = \frac{\mathbf{v}.\mathbf{w}}{||\mathbf{v}||.||\mathbf{w}||}$$

$$\cos \theta = \frac{\text{inner product of vectors } \mathbf{v}, \mathbf{w}}{(\text{length of vector } \mathbf{v}).(\text{length of vector } \mathbf{w})}$$

Example

$$\mathbf{v} = \left(\begin{array}{c} 0 \\ 1 \end{array}\right) \qquad \mathbf{w} = \left(\begin{array}{c} 1 \\ 1 \end{array}\right)$$

$$\mathbf{v}.\,\mathbf{w} = (0.1) + (1.1)$$

$$||\mathbf{v}|| = \sqrt{(0^2 + 1^2)} = 1$$

$$||\mathbf{w}|| = \sqrt{(1^2 + 1^2)} = \sqrt{2}$$

$$\cos \theta = \frac{1}{1.\sqrt{2}} = 0.707$$

$$\theta = \cos^{-1}(0.707) = 45^{\circ}$$

For identical documents

$$v = (1, 1)$$
 and $w = (1, 1)$

$$-$$
 v.w = (1.1) + (1.1) = 2

$$| V | = \sqrt{2}$$

$$\| \mathbf{w} \| = \sqrt{2}$$

$$\bullet \cos \theta = \frac{2}{\sqrt{2}\sqrt{2}} = 1$$

$$\theta = 0^{\circ}$$

For dissimilar documents

$$\mathbf{v} = (1, 0) \text{ and } \mathbf{w} = (0, 1)$$

$$-$$
 v.w = (1.0) + (0.1) = 0

$$\bullet \cos \theta = \frac{0}{\sqrt{1}\sqrt{1}} = 0$$

$$\theta = 90^{\circ}$$

For *n*-dimensional vectors

$$\cos \theta = \frac{\sum_{i=1}^{n} \mathbf{v}_{i} \cdot \mathbf{w}_{i}}{\left(\sum_{i=1}^{n} (\mathbf{v}_{i})^{2}\right)^{1/2} \left(\sum_{i=1}^{n} (\mathbf{w}_{i})^{2}\right)^{1/2}}$$

Example:

- 3-dimensions
- $\mathbf{v} = (v_1, v_2, v_3)$ and $\mathbf{w} = (w_1, w_2, w_3)$

$$\cos \theta = \frac{(\mathbf{v}_1.\mathbf{w}_1) + (\mathbf{v}_2.\mathbf{w}_2) + (\mathbf{v}_3.\mathbf{w}_3)}{\sqrt{[(\mathbf{v}_1)^2 + (\mathbf{v}_2)^2 + (\mathbf{v}_3)^2]} \cdot \sqrt{[(\mathbf{w}_1)^2 + (\mathbf{w}_2)^2 + (\mathbf{w}_3)^2]}}$$

$$\mathbf{d}_1 = \begin{pmatrix} 1 \\ 0 \\ 1 \end{pmatrix} \qquad \mathbf{d}_2 = \begin{pmatrix} 1 \\ 1 \\ 1 \end{pmatrix} \qquad \mathbf{d}_3 = \begin{pmatrix} 0 \\ 1 \\ 1 \end{pmatrix}$$

Example
$$\cos \theta_{12} = \frac{(\mathbf{d}_{11}.\mathbf{d}_{21}) + (\mathbf{d}_{12}.\mathbf{d}_{22}) + (\mathbf{d}_{13}.\mathbf{d}_{23})}{\sqrt{[(\mathbf{d}_{11})^2 + (\mathbf{d}_{12})^2 + (\mathbf{d}_{13})^2]} \cdot \sqrt{[(\mathbf{d}_{21})^2 + (\mathbf{d}_{22})^2 + (\mathbf{d}_{23})^2]}}$$

$$= \frac{(1.1) + (0.1) + (1.1)}{\sqrt{[(1)^2 + (0)^2 + (1)^2]} \cdot \sqrt{[(1)^2 + (1)^2 + (1)^2]}}$$

$$\cos \theta_{12} = \frac{1 + 0 + 1}{\sqrt{2} \cdot \sqrt{3}} = 0.82 \quad (35^\circ)$$

$$\cos \theta = 0.82$$
 0.5 0.82 angle = 35 60 35

- Using one or more queries

$$\mathbf{q}_1 = \left(\begin{array}{c} 0 \\ 0 \\ 1 \end{array}\right) \quad \mathbf{q}_2 = \left(\begin{array}{c} 0 \\ 1 \\ 0 \end{array}\right) \quad \mathbf{q}_3 = \left(\begin{array}{c} 1 \\ 0 \\ 0 \end{array}\right)$$

d1*q1	1	0	1	d2*q1	1	0	1	d3*q1	0	0	1
	0	0	1.414		1	0	1.732		1	0	1.414
	1	1			1	1			1	1	
cosθ =			0.707				0.577				0.707
angle =			45				55				45
d1*q2	1	0	0	d2*q2	1	0	1	d3*q2	0	0	1
	0	1	1.414		1	1	1.732		1	1	1.414
	1	0			1	0			1	0	
cosθ =			0				0.577				0.707
angle =			90				55				45
d1*q3	1	1	1	d2*q3	1	1	1	d3*q3	1	1	1
	0	0	1.414		1	0	1.732		1	0	1.414
	1	0			1	0			0	0	
cos0 =			0.707				0.577				0.707
angle =	:		45				55				45



Example

$$= \begin{pmatrix} 0 \\ 0 \\ 0 \\ 1 \\ 0 \\ 0 \\ 1 \\ 1 \\ 1 \\ 1 \end{pmatrix}$$

$$\mathbf{d}_{3} = \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 1 \end{bmatrix}$$

_				

d1*q1	1	1	1	d2*q1	0	1	0	d3*q1	1	1	1
	1	0	2.828		0	0	2.236		1	0	3.162
	1	0			0	0			1	0	
	0	0			1	0			1	0	
	1	0	0.354		0	0	0		1	0	0.316
	0	0			0	0			1	0	
	1	0			1	0			1	0	
	1	0			1	0			1	0	
	1	0			1	0			1	0	
	1	0			1	0			1	0	
d1*q2	1	1	3	d2*q2	0	1	0	d3*q2	1	1	3
	1	0	4.899		0	0	3.873		1	0	5.477
	1	1			0	1			1	1	
	0	0			1	0			1	0	
	1	1	0.612		0	1	0		1	1	0.548
	0	0			0	0			1	0	
	1	0			1	0			1	0	
	1	0			1	0			1	0	
	1	0			1	0			1	0	
	1	0			1	0			1	0	
d1*q3	1	1	4	d2*q3	0	1	3	d3*q3	1	1	5
	1	0	6.325		0	0	5		1	0	7.071
	1	0			0	0			1	0	
	0	0			1	0			1	0	
	1	0	0.632		0	0	0.6		1	0	0.707
	0	1			0	1			1	1	
	1	1			1	1			1	1	
	1	0			1	0			1	0	
	1	1			1	1			1	1	
	1	1			1	1			1	1	

$$\mathbf{d}_{1} = \begin{vmatrix} 1 \\ 3 \\ 1 \\ 0 \\ 1 \\ 2 \\ \vdots \end{vmatrix}$$

$$\mathbf{d}_{1} = \begin{pmatrix} 1 \\ 2 \\ 1 \\ 3 \\ 1 \\ 0 \\ 1 \\ 2 \\ 1 \\ 1 \end{pmatrix} \qquad \mathbf{d}_{2} = \begin{pmatrix} 0 \\ 4 \\ 0 \\ 1 \\ 0 \\ 0 \\ 1 \\ 2 \\ 1 \\ 1 \end{pmatrix} \qquad \mathbf{d}_{3} = \begin{pmatrix} 2 \\ 1 \\ 1 \\ 0 \\ 0 \\ 1 \\ 2 \\ 0 \\ 1 \\ 2 \end{pmatrix}$$

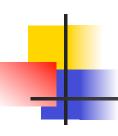
d1*q1	1 2 1 3 1 0 1 2 1	1 0 0 0 0 0 0 0	1 4.796 0.209	d2*q1	0 4 0 1 0 0 1 2 1	1 0 0 0 0 0 0 0	0 4.899 0	d3*q1	2 1 0 1 1 2 0 1 2	1 0 0 0 0 0 0 0	2 4.123 0.485
d1*q2	1 2 1 3 1 0 1 2 1	1 0 1 0 1 0 0 0 0	3 8.307 0.361	d2*q2	0 4 0 1 0 0 1 2 1	1 0 1 0 1 0 0 0 0	0 8.485 0	d3*q2	2 1 0 1 1 2 0 1 2	1 0 1 0 1 0 0 0 0	4 7.141 0.56
d1*q3	1 2 1 3 1 0 1 2 1	1 0 0 0 0 1 1 0 1	4 10.724 0.373	d2*q3	0 4 0 1 0 0 1 2 1	1 0 0 0 0 1 1 0 1	3 10.954 0.274	d3*q3	1 1 0 2 1 1 0 1 2	1 0 0 0 0 1 1 0 1	6 8.367 0.717

Relevance judgment for IR

- A measurement of the outcome of a search or retrieval
- The judgment on what should or should not be retrieved.
- There is no simple answer to what is relevant and what is not relevant: need human users.
 - difficult to define
 - subjective
 - depending on knowledge, needs, time,, etc.
- The central concept of information retrieval

Precision and Recall

- Given a query:
 - Are all retrieved documents relevant?
 - Have all the relevant documents been retrieved?
- Measures for system performance:
 - The first question is about the precision of the search
 - The second is about the completeness (recall) of the search.



Precision and Recall (cont)

	Relevant	Not Relevant
Retrieved	a	b
Not retrieved	С	d

Precision and Recall (cont)

- Precision measures how precise a search is.
 - the higher the precision,
 - the less unwanted documents.

```
Precision = -----

Total number of documents retrieved
```

- Recall measures how complete a search is.
 - the higher the recall,
 - the less missing documents.

Relationship of R and P

- Theoretically,
 - R and P not depend on each other.
- Practically,
 - High Recall is achieved at the expense of precision.
 - High Precision is achieved at the expense of recall.
- When will p = 0?
 - Only when none of the retrieved documents is relevant.
- When will p=1?
 - Only when every retrieved documents are relevant.
- Depending on application, you may want a higher precision or a higher recall.

Alternative measures

Combining recall and precision, F score
 2PR

- Breakeven point: when p = r
- These two measures are commonly used in text mining: classification and clustering.
- Accuracy is not normally used in text domain because the set of relevant documents is almost always very small compared to the set of irrelevant documents.

Web Search as a huge IR system

- A Web crawler (robot) crawls the Web to collect all the pages.
 - Servers establish a huge inverted indexing database and other indexing databases
 - At query (search) time, search engines conduct different types of vector query matching



- The real differences among different search engines are
 - their indexing weight schemes
 - their query process methods
 - their ranking algorithms
 - None of these are published by any of the search engines firms.